



# Intro to Visualization in Python - Static Plots - 1

*One should look for what is and not what he thinks should be. (Albert Einstein)*

# Module completion checklist

Objective	Complete
Prepare data for visualization	
Create histograms, boxplots, and bar charts	

# Visualizing data with matplotlib

- `matplotlib` is a popular plotting library among scientists and data analysts
- It is one of the older Python plotting libraries, so it has become quite flexible and **well-documented** ([link](#))
- Other plotting libraries include Seaborn (which is built on `matplotlib`), `ggplot` (the Python version of the popular R plotting library), `Plotly`, `Bokeh`, and many others
- Pandas also come with some plotting capabilities, and these are just based on `matplotlib`
- Explore the different types of plots that can be created with `matplotlib` by browsing the **gallery** ([link](#))



# Loading packages

- Load the packages:

```
import pandas as pd
import numpy as np
import pickle
import os
from pathlib import Path
```

# Directory settings

- In order to maximize the efficiency of the workflow, encode the directory structure into variables
- Use the `pathlib` library
- The `main_dir` is the variable corresponding to the course folder
- The `data_dir` is the variable corresponding to the data folder

```
# Set 'main_dir' to location of the project folder
home_dir = Path(".").resolve()
main_dir = home_dir.parent.parent
print(main_dir)
```

```
data_dir = str(main_dir) + "/data"
print(data_dir)
```

```
plot_dir = str(main_dir) + "/plots"
if not os.path.exists(plot_dir):
    os.makedirs(plot_dir)
print(plot_dir)
```

# Importing matplotlib

- Import pyplot as plt so that plt.[any\_function]() with appropriate arguments to create a plot
- The pyplot module of the matplotlib library has a large and diverse set of functions
- It allows the ability to create pretty much any conceivable visualization out there
- See the documentation on pyplot [here\(link\)](#)

```
import matplotlib.pyplot as plt
```

## matplotlib.pyplot

matplotlib.pyplot is a state-based interface to matplotlib. It provides a MATLAB-like way of plotting.

pyplot is mainly intended for interactive plots and simple cases of programmatic plot generation:

```
import numpy as np
import matplotlib.pyplot as plt

x = np.arange(0, 5, 0.1)
y = np.sin(x)
plt.plot(x, y)
```

The object-oriented API is recommended for more complex plots.

## Functions

<code>acorr(x, *[, data])</code>	Plot the autocorrelation of x.
<code>angle_spectrum(x[, Fs, Fc, window, pad_to, ...])</code>	Plot the angle spectrum.
<code>annotate(s, xy, *args, **kwargs)</code>	Annotate the point xy with text s.
<code>arrow(x, y, dx, dy, **kwargs)</code>	Add an arrow to the axes.
<code>autoscale([enable, axis, tight])</code>	Autoscale the axis view to the data (toggle).

# Dataset for visualization

- Load the dataset and save it as df

```
# This dataset is of type dataframe. Let's assign this dataset to a variable, so that we can  
manipulate it freely.  
df = pd.read_csv(str(data_dir)+"/" + "diabetes.csv")
```

```
print(type(df))  #<- a Pandas DataFrame!
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
print(len(df))  #<- returns the number of rows
```

```
768
```

# Subsetting data

- Create a subset of the data so that the variables needed are present
- Name this subset as `df_subset`

```
df_subset = df[['DiabetesPedigreeFunction', 'Glucose', 'BloodPressure', 'Age', 'SkinThickness',  
'BMI', 'Insulin', 'Outcome', 'Pregnancies']]  
print(df_subset.head())
```

	DiabetesPedigreeFunction	Glucose	BloodPressure	Age	SkinThickness	BMI	Insulin	
Outcome	Pregnancies							
0		0.627	148	72	50	35	33.6	0
1	6							
1		0.351	85	66	31	29	26.6	0
0	1							
2		0.672	183	64	32	0	23.3	0
1	8							
3		0.167	89	66	21	23	28.1	94
0	1							
4		2.288	137	40	33	35	43.1	168
1	0							

- These variables are chosen because they illustrate the concepts best
- However, you should be able to work with (and visualize) all of your data



# Data reshaping: wide vs. long

- Talking about data reshaping usually refers to converting between what is called either **wide** or **long** data formats
  - **Wide** data is much more visually digestible, which is why it's more common when using data from some type of report
  - **Long** data is much easier to work with in Pandas, and generally speaking in most data analysis and plotting tools

# Data reshaping: wide vs long (cont'd)

- **Wide data** often appears when the values are some type of aggregate (we will use the mean of groups)
- Let's make a typical **wide dataframe** of two rows and eight columns that looks like this:

	Outcome	DiabetesPedigreeFunction	Glucose	BloodPressure	
0	0	0.429734	109.980000	68.184000	31.1
1	1	0.550500	141.257463	70.824627	37.0

# Prepare data: group and summarize

- After grouping and summarizing data, create a summary dataset that will include the following:
  - Grouped data by Target variable
  - Mean value computed on the grouped data that includes the following variables:
    - DiabetesPedigreeFunction
    - Glucose
    - BloodPressure
    - Age
    - SkinThickness
    - BMI
    - Insulin

# Prepare data: group and summarize (cont'd)

- For demonstration, use the original dataframe df to identify the grouping column
- Then use this column to perform the groupby operation and find the mean of the columns present in df\_subset

```
col_dict = df_subset.nunique().to_dict()
grouping_col = min(col_dict, key=col_dict.get)
# Group data by variable with min levels.
grouped = df_subset.groupby(grouping_col)
```

```
# Compute mean on the listed variables using the grouped data.
df_grouped_mean = grouped.mean()[['DiabetesPedigreeFunction', 'Glucose', 'BloodPressure', 'Age',
'SkinThickness', 'BMI', 'Insulin']]
print(df_grouped_mean)
```

	DiabetesPedigreeFunction	Glucose	BloodPressure	Age	SkinThickness
BMI					
Outcome					
0	0.429734	109.980000	68.184000	31.190000	19.664000
30.304200	68.792000				
1	0.550500	141.257463	70.824627	37.067164	22.164179
35.142537	100.335821				

# Prepare data: group and summarize (cont'd)

```
# Reset index of the dataset.  
df_grouped_mean = df_grouped_mean.reset_index()  
print(df_grouped_mean)
```

	Outcome	DiabetesPedigreeFunction	Glucose	BloodPressure	Age	SkinThickness
BMI	Insulin					
0	0	0.429734	109.980000	68.184000	31.190000	19.664000
30.304200	68.792000					
1	1	0.550500	141.257463	70.824627	37.067164	22.164179
35.142537	100.335821					

- This dataframe is considered **wide** because each variable has its own column
- It makes the table easier to present, but inconvenient to run analyses on or visualize

# Why long?

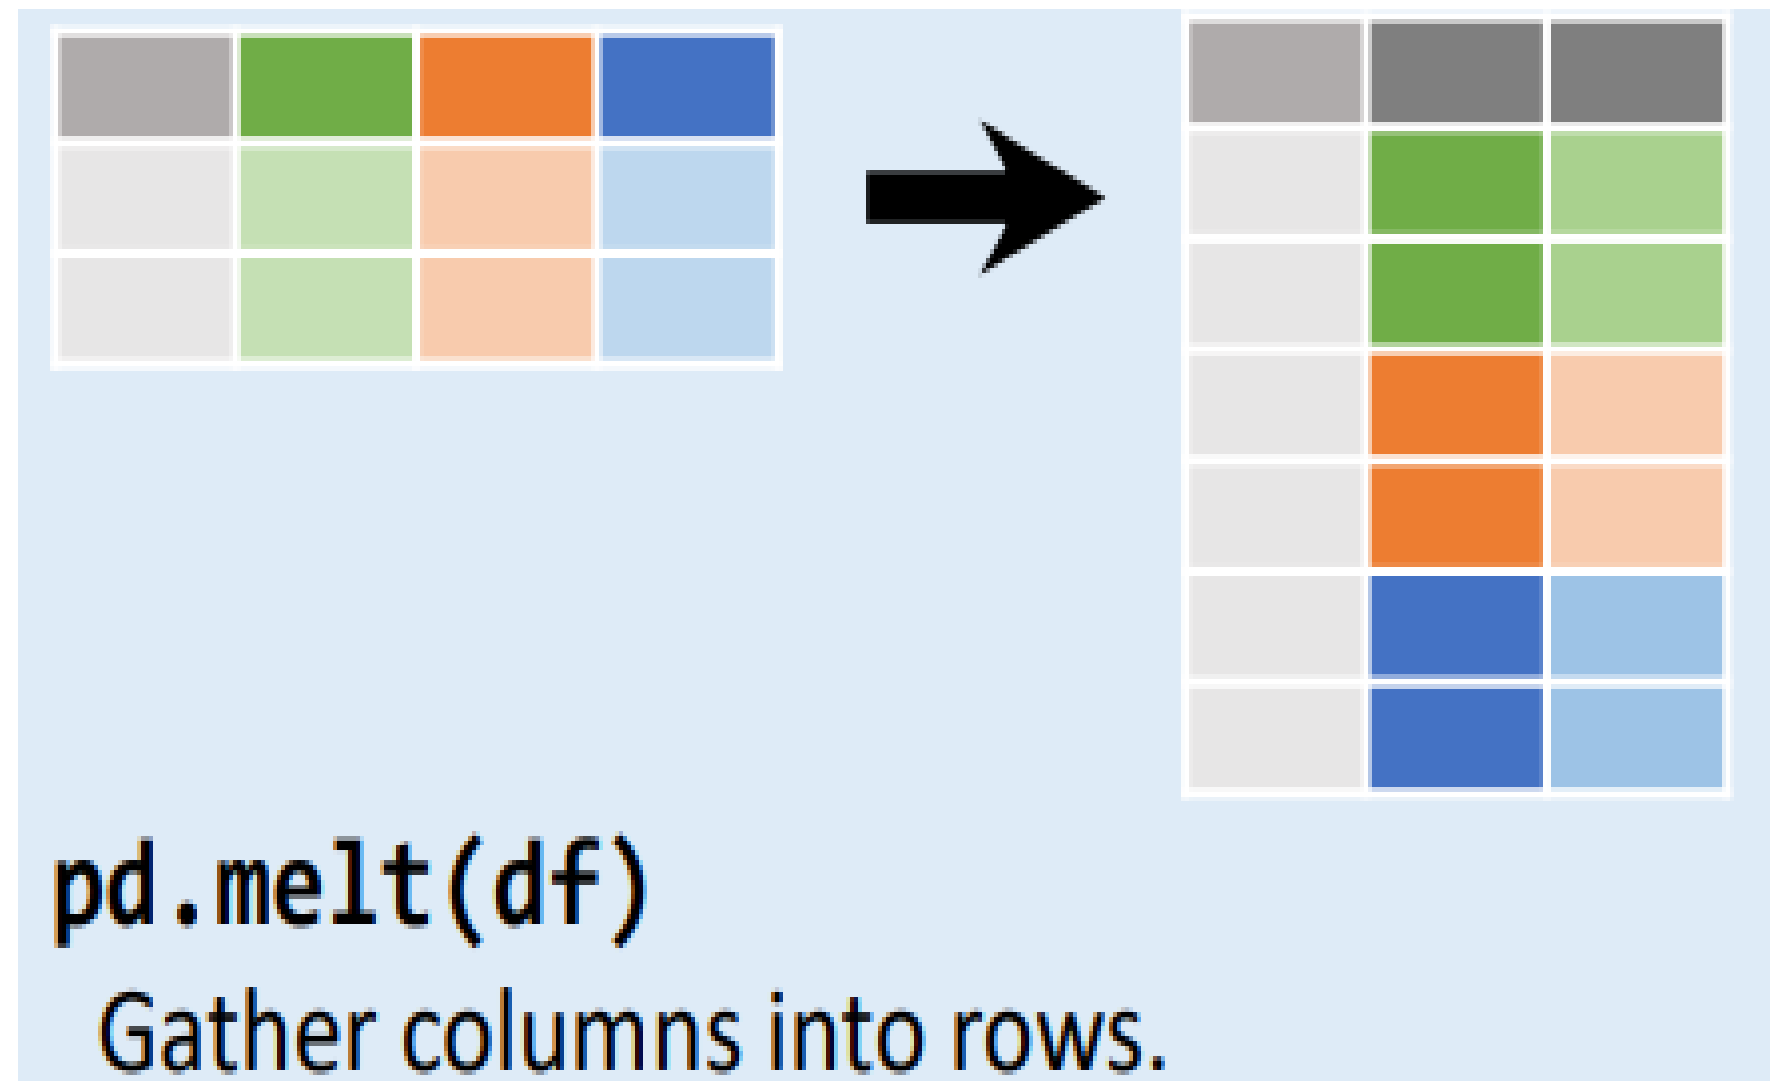
- Now convert this wide data to the **long format**
  - Leave the categorical variable and the mean values as is in their columns
  - All of other variables will appear as a single metric column

- This format is convenient to work with when running analysis and plot the data

	stroke	metric	mean
0	0	avg_glucose_level	104.795513
1	1	avg_glucose_level	132.544739
2	0	bmi	28.825388
3	1	bmi	30.217789

# Wide to long format: melt

- To **convert from wide to long format**, use the Pandas `melt` function with the following arguments:
  - i. Wide dataframe
  - ii. Variable(s) that will be preserved as the `ids` of the data (like categorical variables)
  - iii. Name of the variable that will now contain the column names from the wide data selected to melt together
  - iv. Name of the column that will contain respective values corresponding to the melted columns



# Wide to long format: melt (cont'd)

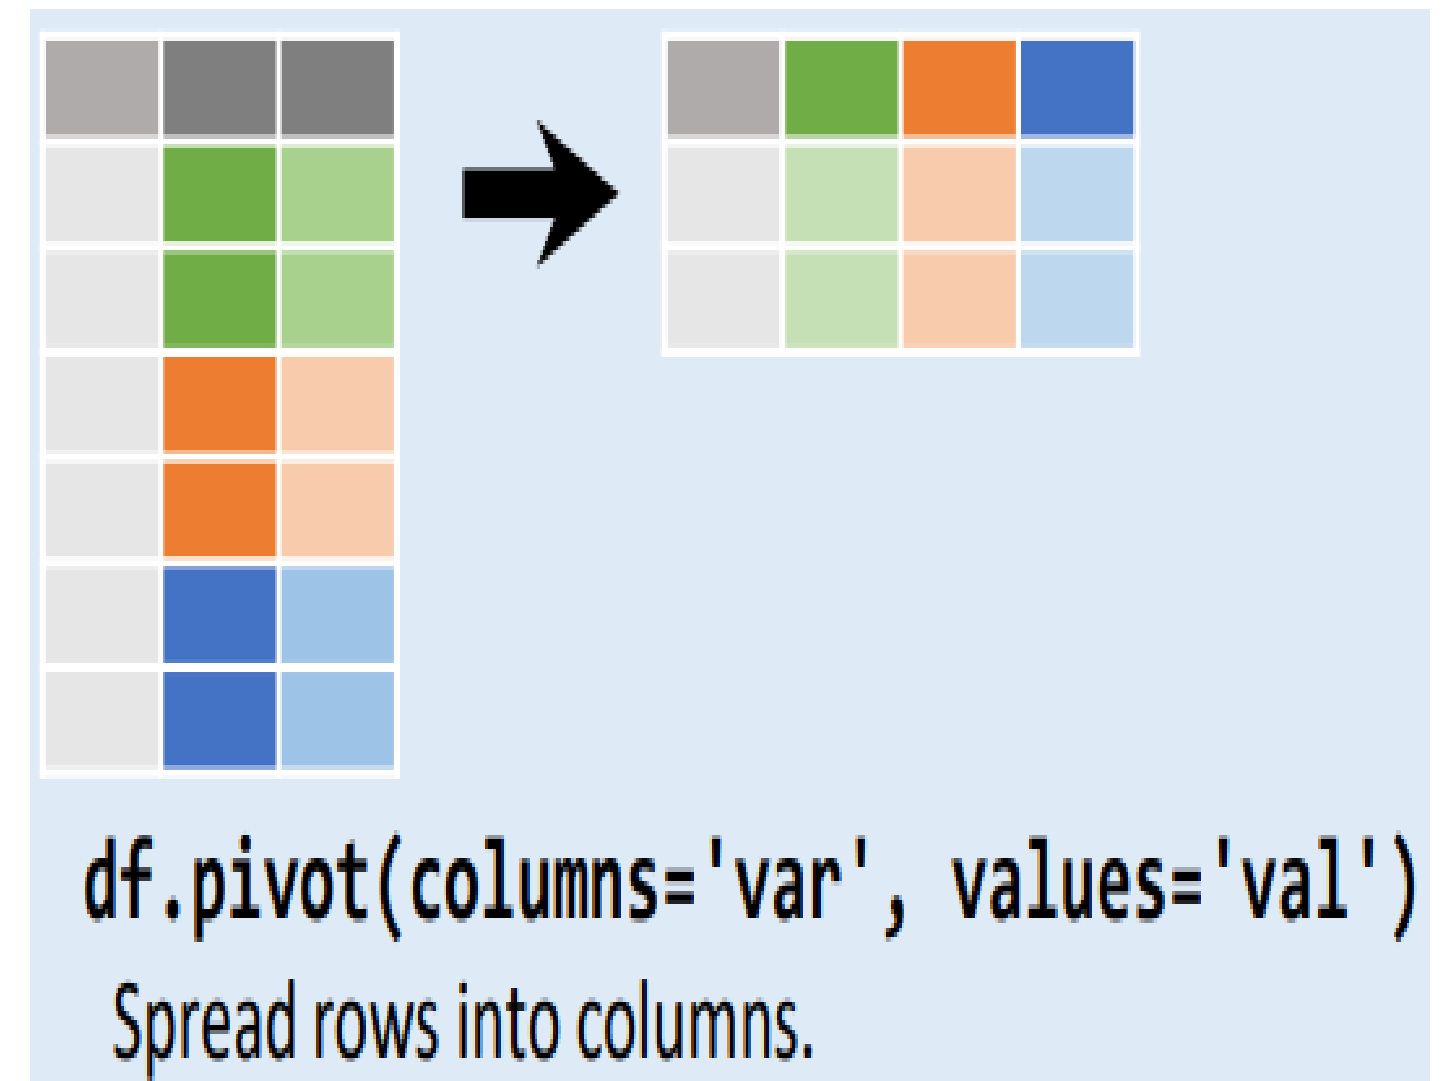
```
# Melt the wide data into long.
df_grouped_mean_long = pd.melt(df_grouped_mean,          #<- wide dataset
                               id_vars = [grouping_col],  #<- identifying variable
                               var_name = 'metric',        #<- contains col names of wide data
                               value_name = 'mean')        #<- contains values from above columns
print(df_grouped_mean_long)
```

	Outcome		metric	mean
0	0	DiabetesPedigreeFunction		0.429734
1	1	DiabetesPedigreeFunction		0.550500
2	0	Glucose		109.980000
3	1	Glucose		141.257463
4	0	BloodPressure		68.184000
5	1	BloodPressure		70.824627
6	0	Age		31.190000
7	1	Age		37.067164
8	0	SkinThickness		19.664000
9	1	SkinThickness		22.164179
10	0	BMI		30.304200
11	1	BMI		35.142537
12	0	Insulin		68.792000
13	1	Insulin		100.335821



# Long to wide format: pivot

- Convert the **long data back to wide** format with the `.pivot()` method
1. The `index` argument refers to what values will become the `ids` in the new dataframe
  2. The `columns` argument refers to the column in which its values will be converted to column names
  3. Lastly, supply the `values` argument to fill in the values of the wide data



# Long to wide format: pivot (cont'd)

```
# Melt the long data into wide.
df_grouped_mean_wide = df_grouped_mean_long.pivot(
    index = [grouping_col],    #<- identifying
    variable = ...,            #<- variable names of long data
    columns = 'metric',        #<- col names of wide data
    values = 'mean')           #<- values from above
print(df_grouped_mean_wide)
```

	metric	Age	BMI	BloodPressure	DiabetesPedigreeFunction	Glucose
Outcome	Insulin	SkinThickness				
0	31.190000	30.304200	68.184000	0.429734	109.980000	
1	37.067164	35.142537	70.824627	0.550500	141.257463	

# Module completion checklist

Objective	Complete
Prepare data for visualization	✓
Create histograms, boxplots, and bar charts	


# Univariate plots

- Univariate plots are used to **visualize the distribution of a single variable**
- They are mainly used in the initial stages of EDA when we want to learn more about individual variables in our data
- They are also combined with other univariate plots to compare data distributions of different variables
- Univariate plots include the following popular graphs: histogram, boxplot, density curve, dot plot, QQ plot, and bar plot

# Univariate plots: histogram

- A histogram represents the **distribution of numerical data**
- The height of each bar has been calculated as the number of observations in that range
- Use `plt.hist()` to produce a basic histogram of any numeric variable


```
plt.rcParams.update({'font.size': 15})  
plt.hist(df_subset['DiabetesPedigreeFunction'])  
plt.show()
```



# Univariate plots: histogram (cont'd)

- Bins represent the intervals in which to group the observations
- Control the number of bins with the `bins` parameter
- As the **number of bins increases**, the range of values each bin represents decreases, and so does the height of the bar


```
plt.hist(df_subset['DiabetesPedigreeFunction'],  
bins = 20)  
plt.xlabel('DiabetesPedigreeFunction')      #<-  
label x-axis  
plt.ylabel('Frequency')                    #<- label y-axis  
plt.title('DiabetesPedigreeFunction  
distribution')                             #<- add plot title  
plt.show()
```



# Univariate plots: boxplot

- A boxplot is a visual summary of the **25th, 50th, and 75th percentiles**
- The orange line shows the median of `ppl_total`
- The top and bottom of the box are the 25th and 75th percentile respectively
- The outermost lines are called the **whiskers**
- Values beyond whiskers are considered **outliers** - they are substantially outside the rest of the data

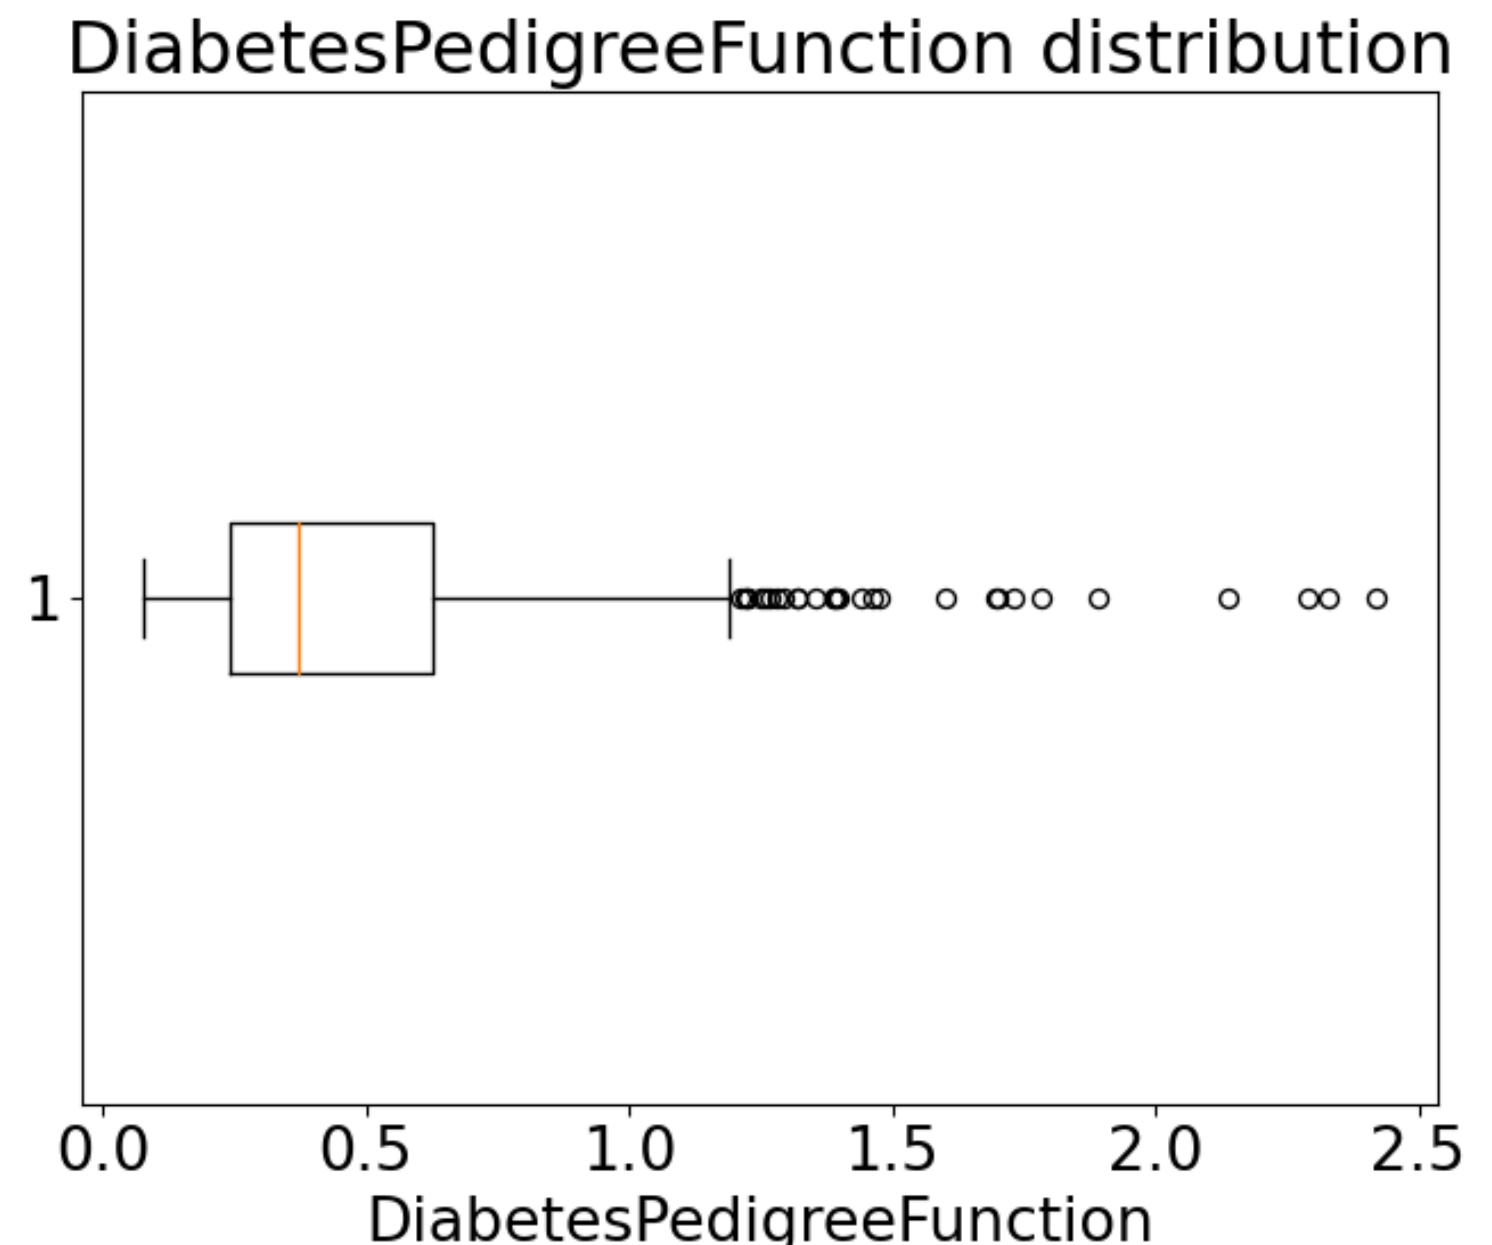
```
plt.boxplot(df_subset['DiabetesPedigreeFunction'])  
plt.show()
```



# Univariate plots: boxplot (cont'd)

- The orientation of the plot can be set to horizontal by setting `vert = False`
- Answer in chat: By looking at this boxplot, what can be told about the **'DiabetesPedigreeFunction'** distribution in the data?

```
plt.boxplot(df_subset['DiabetesPedigreeFunction'],  
            vert = False)  
plt.xlabel('DiabetesPedigreeFunction')      #  
            label x-axis  
plt.title('DiabetesPedigreeFunction  
distribution')      # add plot title  
plt.show()
```





# Univariate plots: bar chart

- A bar chart is a plot where the height of each bar represents **the numeric value of a category**
- Use `plt.bar()` to produce a basic histogram of **any categorical variable**
- Bar charts are most commonly used when visualizing survey data or summary data
- The general syntax for creating a bar chart consists of 3 main variables:
  - position of the bars on the `axis`
  - height of the bars
  - names of categories that are used to label the bars

# Univariate plots: bar chart - cont'd

- When plotting bar charts of any complexity, the best type of data to use is **long data**
- First create a simple bar chart of the variable means using the `df_grouped_mean_long` data we created earlier

```
print(df_grouped_mean_long.head())
```

# Univariate plots: bar chart - cont'd

- Next, filter 'Outcome' by a category and only keep two columns: metric and mean

```
query = 'Outcome' + "==" + str('0')  
df_true_means = df_grouped_mean_long.query(query)[['metric', 'mean']]  
print(df_true_means)
```

# Univariate plots: bar chart - cont'd

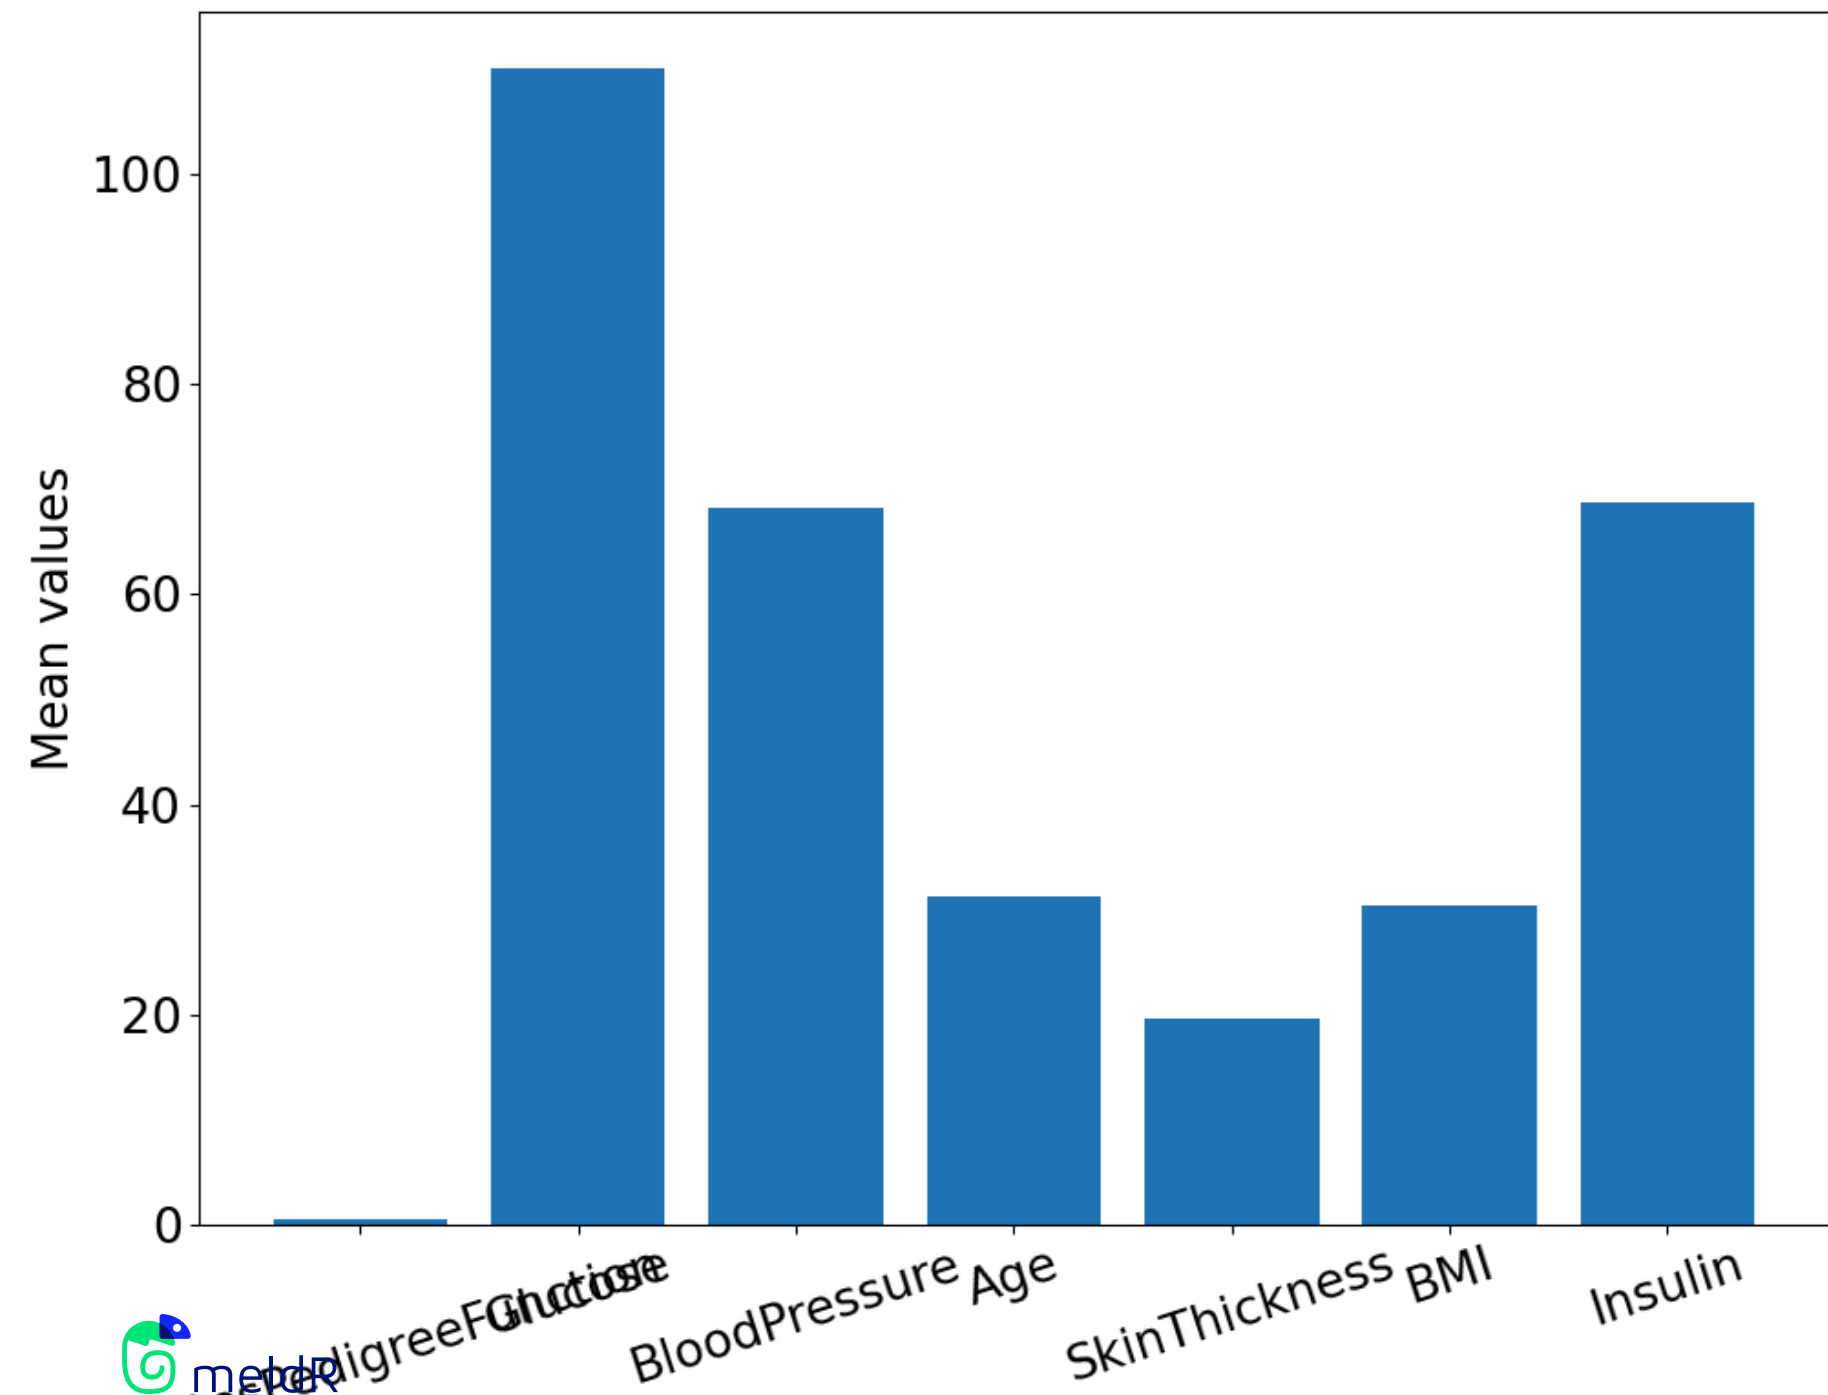
- Now, get the data needed and assign it to the three variables for convenience and clarity:
  1. **Categories** (i.e., labels) that will represent each bar are all contained in the `metric` column
  2. **Bar heights** are contained in the `mean` column for each of the 5 categories
  3. **Bar positions** will be a range of numbers based on the number of categories (i.e., bars)

```
bar_labels = df_true_means['metric']      #<- 1  
bar_heights = df_true_means['mean']      #<- 2  
num_bars = len(bar_heights)  
bar_positions = np.arange(num_bars)      #<- 3
```

# Univariate plots: bar chart - cont'd

- Labels are tricky to fit sometimes, so either **adjust** the figure size or label orientation

```
plt.figure(figsize = (12, 9))  
plt.bar(bar_positions, bar_heights)  
plt.xticks(bar_positions, bar_labels, rotation = 18)  
plt.ylabel('Mean values')
```



# Customize anything

- All possible style customizations are available in a `matplotlibrc` file
- **This sample (link)** contains all of them, and any of those parameters can be passed to `rcParams` variable like we did earlier
- This sample contains a script of parameters and their default values
- Here's a part of that file with a sample of all parameters for modifying the style of the axes

```
## *****  
## * AXES *  
## *****  
## Following are default face and edge colors, default tick sizes,  
## default font sizes for tick labels, and so on. See  
## https://matplotlib.org/api/axes\_api.html#module-matplotlib.axes  
#axes.facecolor:      white      # axes background color  
#axes.edgecolor:      black      # axes edge color  
#axes.linewidth:      0.8        # edge line width  
#axes.grid:           False      # display grid or not  
#axes.grid.axis:      both       # which axis the grid should apply to  
#axes.grid.which:      major      # grid lines at {major, minor, both} ticks  
#axes.titlelocation:  center     # alignment of the title: {left, right, center}
```

# Knowledge check



# Module completion checklist

Objective	Complete
Prepare data for visualization	✓
Create histograms, boxplots, and bar charts	✓



# Congratulations on completing this module!

You are now ready to try Tasks 1-13 in the Exercise for this topic

